

Flipping The Dialogue: Training and Evaluating User Language Models



Tarek Naous



Philippe Laban



Wei Xu



Jennifer Neville



<https://huggingface.co/microsoft/UserLM-8b>

You must act as a real user and make requests as a user would to an AI assistant. You should be brief and potentially make [...]

The task you are trying to accomplish as a user is: **getting information about strategies for losing weight** and the impact of certain drugs or medications on weight gain.

You must act as a real user and make requests as a user would to an AI assistant. You should be brief and potentially make [...]

The task you are trying to accomplish as a user is: **getting information about strategies for losing weight** and the impact of certain drugs or medications on weight gain.



hey, looking for tips on losing weight + how certain meds/drugs affect weight gain (like antidepressants, birth control, etc). any info?

You must act as a real user and make requests as a user would to an AI assistant. You should be brief and potentially make [...]

The task you are trying to accomplish as a user is: **getting information about strategies for losing weight** and the impact of certain drugs or medications on weight gain.



hey, looking for tips on losing weight + how certain meds/drugs affect weight gain (like antidepressants, birth control, etc). any info?



hey can you tell me about ozempic and other weight loss drugs? also what are some good strategies for losing weight in general



hey can u help me out? i've been trying to lose some weight but i feel like im stuck. also i started some new meds recently and i swear im gaining weight faster now.

You must act as a real user and make requests as a user would to an AI assistant. You should be brief and potentially make [...]

The task you are trying to accomplish as a user is: **getting information about strategies for losing weight** and the impact of certain drugs or medications on weight gain.



hey, looking for tips on losing weight + how certain meds/drugs affect weight gain (like antidepressants, birth control, etc). any info?



hey can you tell me about ozempic and other weight loss drugs? also what are some good strategies for losing weight in general



hey can u help me out? i've been trying to lose some weight but i feel like im stuck. also i started some new meds recently and i swear im gaining weight faster now.

Real User



You must act as a real user and make requests as a user would to an AI assistant. You should be brief and potentially make [...]

The task you are trying to accomplish as a user is: **getting information about strategies for losing weight** and the impact of certain drugs or medications on weight gain.



hey, looking for tips on losing weight + how certain meds/drugs affect weight gain (like antidepressants, birth control, etc). any info?



hey can you tell me about ozempic and other weight loss drugs? also what are some good strategies for losing weight in general



hey can u help me out? i've been trying to lose some weight but i feel like im stuck. also i started some new meds recently and i swear im gaining weight faster now.

Real User



does allertine cause weight gain

You must act as a real user and make requests as a user would to an AI assistant. You should be brief and potentially make [...]

The task you are trying to accomplish as a user is: **getting information about strategies for losing weight** and the impact of certain drugs or medications on weight gain.



hey, looking for tips on losing weight + how certain meds/drugs affect weight gain (like antidepressants, birth control, etc). any info?



hey can you tell me about ozempic and other weight loss drugs? also what are some good strategies for losing weight in general



hey can u help me out? i've been trying to lose some weight but i feel like im stuck. also i started some new meds recently and i swear im gaining weight faster now.

Real User



does allertine cause weight gain

The inherent assistant role of language models makes them poor user simulators

SIMULATING USERS IN CONVERSATIONS ...

User Intent: Write a Python function: given an array of integers, sort ones between 1 and 9 inclusive, reverse the array, and replace digits by their name from "One", "Two", "Three", etc.

SIMULATING USERS IN CONVERSATIONS ...

... USING AN ASSISTANT LANGUAGE MODEL

User Intent: Write a Python function: given an array of integers, sort ones between 1 and 9 inclusive, reverse the array, and replace digits by their name from "One", "Two", "Three", etc.

 **GPT-4o**

-  Turn digits into names in a list
-  ... 
-  I want to sort the numbers if they're between 1 and 9
-  ... 
-  I want to flip the list after that
-  ... 
-  I want to use names for numbers from One to Nine

**Simple & Direct
User Turns**



✓ Success



Benchmark Score: 74.6 ↑

SIMULATING USERS IN CONVERSATIONS ...

... USING AN ASSISTANT LANGUAGE MODEL

User Intent: Write a Python function: given an array of integers, sort ones between 1 and 9 inclusive, reverse the array, and replace digits by their name from "One", "Two", "Three", etc.

 **GPT-4o**

-  Turn digits into names in a list
- ...
-  I want to sort the numbers if they're between 1 and 9
- ...
-  I want to flip the list after that
- ...
-  I want to use names for numbers from One to Nine

**Simple & Direct
User Turns**

→  **Success** 

 **Benchmark Score:** 74.6 

... USING A USER LANGUAGE MODEL

 **UserLM-8b**

-  ignore any number that is not in the 1-to-9 range
- ...
-  now add sorting into it
- ...
-  now reverse the sorted list
- ...
-  translate each remaining digit into its capitalized English name

**Nuanced & Indirect
User Turns**

→  **Failure** 

 **Benchmark Score:** 57.4 

SIMULATING USERS IN CONVERSATIONS ...

... USING AN ASSISTANT LANGUAGE MODEL

... USING A USER LANGUAGE MODEL

User Intent: Write a Python function: given an array of integers, sort ones between 1 and 9 inclusive, reverse the array, and replace digits by their name from "One", "Two", "Three", etc.

GPT-4o

Turn digits into names in a list

...

I want to sort the numbers if they're between 1 and 9

...

I want to flip the list after that

...

I want to use names for numbers from One to Nine

Simple & Direct User Turns

Success

Benchmark Score: 74.6

Simulator Capability

GPT-4o (Prompted)	Real Users	UserLM-8b (Trained)
Diversity (↑)		
74.4	94.01	94.5
Dialogue Termination (↑)		
1.3	100	63.5
Intent Decomposition (↓)		
7.7	1.68	2.7
Naturalness (↑)		
3.3	90.15	80.2
Robustness (↑)		
54.9	100	94.3

UserLM-8b

ignore any number that is not in the 1-to-9 range

...

now add sorting into it

...

now reverse the sorted list

...

translate each remaining digit into its capitalized English name

Nuanced & Indirect User Turns

X Failure

Benchmark Score: 57.4

Flipping The Dialogue

1

Training User LMs

How do we train a user language model?

User-Assistant Conversations

In-the-wild user-assistant conversations

CONVERSATION 1



does allertine cause weight gain

As an AI language model, I don't have personal experiences. But according to ..



do antihistamines cause weight gain

Most antihistamines do not cause weight gain as a side effect. However ..



how much weight can i lose by cutting sugar for 40 days

The amount of weight you can lose ...



CONVERSATION N-1

...

CONVERSATION N



Is einstein determinism challenged by quantum mechanics?

Quantum mechanics challenges Einstein's determinism, which held that ...



What did Hawkins say?

It seems you are referring to Stephen Hawking, the ...



User-Assistant Conversations

In-the-wild user-assistant conversations

CONVERSATION 1



does allertine cause weight gain

As an AI language model, I don't have personal experiences. But according to ..



do antihistamines cause weight gain

Most antihistamines do not cause weight gain as a side effect. However ..



how much weight can i lose by cutting sugar for 40 days

The amount of weight you can lose ...



CONVERSATION N-1

...

CONVERSATION N



Is einstein determinism challenged by quantum mechanics?

Quantum mechanics challenges Einstein's determinism, which held that ...



What did Hawkins say?

It seems you are referring to Stephen Hawking, the ...



User Intent Generation

Generate generic user intents for conversations

INTENT 1

You are a user talking with an assistant language model to get information about strategies for losing weight and the impact of certain drugs or medications on weight gain.

INTENT N-1

You are a user talking with an assistant language model to ...

INTENT N

You are a user talking with an assistant language model to understand how quantum Mechanics challenges Einstein's determinism and get the perspective of other scientists.

User-Assistant Conversations

In-the-wild user-assistant conversations

CONVERSATION 1



does allertine cause weight gain

As an AI language model, I don't have personal experiences. But according to ..



do antihistamines cause weight gain

Most antihistamines do not cause weight gain as a side effect. However ..



how much weight can i lose by cutting sugar for 40 days

The amount of weight you can lose ...



CONVERSATION N-1

...

CONVERSATION N



Is einstein determinism challenged by quantum mechanics?

Quantum mechanics challenges Einstein's determinism, which held that ...



What did Hawkins say?

It seems you are referring to Stephen Hawking, the ...



User Intent Generation

Generate generic user intents for conversations

1 conversation (3 turns) + intent

yields 4 samples

INTENT 1

You are a user talking with an assistant language model to get information about strategies for losing weight and the impact of certain drugs or medications on weight gain.

INTENT N-1

You are a user talking with an assistant language model to ...

INTENT N

You are a user talking with an assistant language model to understand how quantum Mechanics challenges Einstein's determinism and get the perspective of other scientists.

User Language Modeling

Flip the dialogue turns to create training samples for a User LM

User-Assistant Conversations

In-the-wild user-assistant conversations

CONVERSATION 1



does allertine cause weight gain

As an AI language model, I don't have personal experiences. But according to ..



do antihistamines cause weight gain

Most antihistamines do not cause weight gain as a side effect. However ..



how much weight can i lose by cutting sugar for 40 days

The amount of weight you can lose ...



CONVERSATION N-1

...

CONVERSATION N



Is einstein determinism challenged by quantum mechanics?

Quantum mechanics challenges Einstein's determinism, which held that ...



What did Hawkins say?

It seems you are referring to Stephen Hawking, the ...



User Intent Generation

Generate generic user intents for conversations

1 conversation (3 turns) + intent
yields 4 samples

INTENT 1

You are a user talking with an assistant language model to get information about strategies for losing weight and the impact of certain drugs or medications on weight gain.

INTENT N-1

You are a user talking with an assistant language model to ...

INTENT N

You are a user talking with an assistant language model to understand how quantum Mechanics challenges Einstein's determinism and get the perspective of other scientists.

User Language Modeling

Flip the dialogue turns to create training samples for a User LM

SAMPLE 1

Turn 1

Conversation
Initiation

You are a user chatting with an assistant language model to get information about strategies for losing weight and the impact of certain drugs or medications on weight gain.



OUTPUT: does allertine cause weight gain

User-Assistant Conversations

In-the-wild user-assistant conversations

CONVERSATION 1



does allertine cause weight gain

As an AI language model, I don't have personal experiences. But according to ..





do antihistamines cause weight gain

Most antihistamines do not cause weight gain as a side effect. However ..





how much weight can i lose by cutting sugar for 40 days

The amount of weight you can lose ...



CONVERSATION N-1

...

CONVERSATION N



Is einstein determinism challenged by quantum mechanics?

Quantum mechanics challenges Einstein's determinism, which held that ...






What did Hawkins say?

It seems you are referring to Stephen Hawking, the ...



User Intent Generation

Generate generic user intents for conversations



User Language Modeling

Flip the dialogue turns to create training samples for a User LM

SAMPLE 1

Turn 1

Conversation Initiation



You are a user chatting with an assistant language model to get information about strategies for losing weight and the impact of certain drugs or medications on weight gain.



OUTPUT: does allertine cause weight gain

SAMPLE 2

Turn 2

Interaction



You are a user chatting with an assistant language model to get information about strategies for losing ...



does allertine cause weight gain

As an AI language model, I don't have personal experiences. But according to ...





OUTPUT: do antihistamines cause weight gain

...

User-Assistant Conversations

In-the-wild user-assistant conversations

CONVERSATION 1



does allertine cause weight gain

As an AI language model, I don't have personal experiences. But according to ..



do antihistamines cause weight gain

Most antihistamines do not cause weight gain as a side effect. However ..



how much weight can i lose by cutting sugar for 40 days

The amount of weight you can lose ...



CONVERSATION N-1

...

CONVERSATION N



Is einstein determinism challenged by quantum mechanics?

Quantum mechanics challenges Einstein's determinism, which held that ...



What did Hawkins say?

It seems you are referring to Stephen Hawking, the ...



User Intent Generation

Generate generic user intents for conversations

1 conversation (3 turns) + intent
yields 4 samples

INTENT 1

You are a user talking with an assistant language model to get information about strategies for losing weight and the impact of certain drugs or medications on weight gain.

INTENT N-1

You are a user talking with an assistant language model to ...

INTENT N

You are a user talking with an assistant language model to understand how quantum Mechanics challenges Einstein's determinism and get the perspective of other scientists.

User Language Modeling

Flip the dialogue turns to create training samples for a User LM

SAMPLE 1

Turn 1

*Conversation
Initiation*



You are a user chatting with an assistant language model to get information about strategies for losing weight and the impact of certain drugs or medications on weight gain.



OUTPUT: does allertine cause weight gain

SAMPLE 2

Turn 2

Interaction



You are a user chatting with an assistant language model to get information about strategies for losing ...



does allertine cause weight gain

As an AI language model, I don't have personal experiences. But according to ...



OUTPUT: do antihistamines cause weight gain

SAMPLE 3

Turn 3

Interaction

...

User-Assistant Conversations

In-the-wild user-assistant conversations

CONVERSATION 1



does allertine cause weight gain

As an AI language model, I don't have personal experiences. But according to ..



do antihistamines cause weight gain

Most antihistamines do not cause weight gain as a side effect. However ..



how much weight can i lose by cutting sugar for 40 days

The amount of weight you can lose ...



CONVERSATION N-1

...

CONVERSATION N



Is einstein determinism challenged by quantum mechanics?

Quantum mechanics challenges Einstein's determinism, which held that ...



What did Hawkins say?

It seems you are referring to Stephen Hawking, the ...



User Intent Generation

Generate generic user intents for conversations

1 conversation (3 turns) + intent
yields 4 samples

INTENT 1

You are a user talking with an assistant language model to get information about strategies for losing weight and the impact of certain drugs or medications on weight gain.

INTENT N-1

You are a user talking with an assistant language model to ...

INTENT N

You are a user talking with an assistant language model to understand how quantum Mechanics challenges Einstein's determinism and get the perspective of other scientists.

User Language Modeling

Flip the dialogue turns to create training samples for a User LM

SAMPLE 1

Turn 1

Conversation
Initiation



You are a user chatting with an assistant language model to get information about strategies for losing weight and the impact of certain drugs or medications on weight gain.



OUTPUT: does allertine cause weight gain

SAMPLE 2

Turn 2

Interaction



You are a user chatting with an assistant language model to get information about strategies for losing ...



does allertine cause weight gain

As an AI language model, I don't have personal experiences. But according to ...



OUTPUT: do antihistamines cause weight gain

SAMPLE 3

Turn 3

Interaction

...

SAMPLE 4

Turn K

Conversation
Ending



You are a user chatting with an assistant language model to get information about strategies for losing ...



how much weight can i lose by cutting sugar for 40 days

The amount of weight you can lose ...



OUTPUT: <|endconversation|>

Training Setup and Details

Data: 384,336 conversations from WildChat after filtering/deduplication (1,047,930 training samples)

Intent Generation: generated for each conversation using few-shot prompting with GPT-4o

Models: We train UserLM-1b and UserLM-8b starting from Llama3-8b-Base and Llama3.2-1b-Base

Training Details: max seq length of 2048 tokens, batch size of 1024 samples, learning rate of $2e-5$

Distributional Alignment with Human Utterances

Perplexity on user utterances in two setups:

	WildChat
User Simulator	
Llama3.2-1b-Instruct	
Llama3-8b-Instruct	
GPT-4o-mini	
GPT-4o	
USP-8b	
UserLM-1b	
UserLM-8b	

Distributional Alignment with Human Utterances

Perplexity on user utterances in two setups:

 Without intent conditioning

User Simulator	WildChat
	 (↓)
Llama3.2-1b-Instruct	37.68
Llama3-8b-Instruct	98.29
GPT-4o-mini	26.08
GPT-4o	26.19
USP-8b	32.08
UserLM-1b	8.30
UserLM-8b	5.60

Distributional Alignment with Human Utterances


Perplexity on user utterances in two setups:

 Without intent conditioning

 With intent conditioning

User Simulator	WildChat	
	 (↓)	 (↓)
Llama3.2-1b-Instruct	37.68	29.09
Llama3-8b-Instruct	98.29	48.13
GPT-4o-mini	26.08	16.08
GPT-4o	26.19	21.40
USP-8b	32.08	21.78
UserLM-1b	8.30	7.78
UserLM-8b	5.60	4.33

Distributional Alignment with Human Utterances


		in-domain		out-of-domain	
		WildChat		PRISM	
Perplexity on user utterances in two setups:					
 Without intent conditioning	User Simulator	 (↓)	 (↓)	 (↓)	 (↓)
	Llama3.2-1b-Instruct	37.68	29.09	84.00	53.54
	Llama3-8b-Instruct	98.29	48.13	89.98	40.86
	GPT-4o-mini	26.08	16.08	35.02	20.80
 With intent conditioning	GPT-4o	26.19	21.40	40.25	36.29
	USP-8b	32.08	21.78	50.91	30.16
	UserLM-1b	8.30	7.78	18.58	10.33
	UserLM-8b	5.60	4.33	14.92	7.42

Distributional Alignment with Human Utterances

		in-domain		out-of-domain	
		WildChat		PRISM	
Perplexity on user utterances in two setups:	User Simulator	 (↓)	 (↓)	 (↓)	 (↓)
User LMs show better distributional alignment with user text <i>(better with intent conditioning, generalize to out-of-domain data)</i>					
 With intent conditioning	GPT-4o	26.19	21.40	40.25	36.29
	USP-8b	32.08	21.78	50.91	30.16
	UserLM-1b	8.30	7.78	18.58	10.33
	UserLM-8b	5.60	4.33	14.92	7.42

How important is conditioning on generic intent?


We train User LMs without any first-turn intent conditioning and compare their test-time performance



The biggest PPL gains are observed when models are trained with intent conditioning

Is starting from a base or instruction-tuned checkpoint better?

We compare our User LMs to ones we train starting from the instruct Llama models



Starting from a base checkpoint results in better PPL than ones tuned for the assistant role

Flipping The Dialogue

1

Training User LMs

How do we train a user language model?

2

Intrinsic Evaluation

How do we evaluate a user language model?

Evaluating Multi-turn Interaction Properties

User Simulator

Llama3.2-1b-Instruct

Llama3-8b-Instruct

GPT-4o-mini

GPT-4o

USP-8b

UserLM-1b

UserLM-8b

Human (*estimate*)

Evaluating Multi-turn Interaction Properties

 **First-turn Diversity**
Are first turn requests generated diverse?

User Simulator	 (↑)
Llama3.2-1b-Instruct	81.36
Llama3-8b-Instruct	81.31
GPT-4o-mini	66.10
GPT-4o	74.42
USP-8b	94.37
UserLM-1b	90.90
UserLM-8b	94.55
Human (<i>estimate</i>)	94.01

Evaluating Multi-turn Interaction Properties

 **First-turn Diversity**
Are first turn requests generated diverse?

 **Intent Decomposition**
Does the simulator split its intent across turns?

User Simulator	 (↑)	 (↓)
Llama3.2-1b-Instruct	81.36	15.72
Llama3-8b-Instruct	81.31	23.95
GPT-4o-mini	66.10	9.66
GPT-4o	74.42	7.68
USP-8b	94.37	6.33
UserLM-1b	90.90	3.07
UserLM-8b	94.55	2.69
Human (<i>estimate</i>)	94.01	1.68

Evaluating Multi-turn Interaction Properties

 **First-turn Diversity**
Are first turn requests generated diverse?

 **Intent Decomposition**
Does the simulator split its intent across turns?

 **Dialogue Termination**
Does the simulator at ending the conversation ?

User Simulator	 (↑)	 (↓)	 (↑)
Llama3.2-1b-Instruct	81.36	15.72	3.47
Llama3-8b-Instruct	81.31	23.95	3.51
GPT-4o-mini	66.10	9.66	15.31
GPT-4o	74.42	7.68	1.38
USP-8b	94.37	6.33	21.31
UserLM-1b	90.90	3.07	56.83
UserLM-8b	94.55	2.69	63.54
Human (<i>estimate</i>)	94.01	1.68	—

Evaluating Multi-turn Interaction Properties



First-turn Diversity

Are first turn requests generated diverse?

User Simulator



(↑)



(↓)



(↑)

Llama3.2-1b-Instruct

81.36

15.72

3.47

Llama3.8b-Instruct

81.31

23.95

3.51

User LMs align better with real users properties

(requests are diverse, intents are split across turns, dialogues end)



Dialogue Termination

Does the simulator at ending the conversation ?

UserLM-1b

90.90

3.07

56.83

UserLM-8b

94.55

2.69

63.54

Human (estimate)

94.01

1.68

—

Evaluating Simulator Robustness

User Simulator

Llama3.2-1b-Instruct

Llama3-8b-Instruct

GPT-4o-mini

GPT-4o

USP-8b

UserLM-1b

UserLM-8b

Human (*estimate*)

Evaluating Simulator Robustness



Naturalness

Do simulator utterances resemble natural user text?

User Simulator	 (↑)
Llama3.2-1b-Instruct	0.14
Llama3-8b-Instruct	0.20
GPT-4o-mini	0.04
GPT-4o	3.31
USP-8b	77.73
UserLM-1b	78.96
UserLM-8b	80.21
Human (<i>estimate</i>)	90.15

Evaluating Simulator Robustness



Naturalness

Do simulator utterances resemble natural user text?



User Role Adherence

Does the simulator stick to its role as the user?

User Simulator	 (↑)	 (↑)
Llama3.2-1b-Instruct	0.14	77.55
Llama3-8b-Instruct	0.20	63.25
GPT-4o-mini	0.04	80.20
GPT-4o	3.31	38.85
USP-8b	77.73	98.05
UserLM-1b	78.96	91.30
UserLM-8b	80.21	93.95
Human (<i>estimate</i>)	90.15	—

Evaluating Simulator Robustness



Naturalness

Do simulator utterances resemble natural user text?



User Role Adherence

Does the simulator stick to its role as the user?



User Intent Adherence

Does the simulator stick to its user intent?

User Simulator	 (↑)	 (↑)	 (↑)
Llama3.2-1b-Instruct	0.14	77.55	54.95
Llama3-8b-Instruct	0.20	63.25	78.05
GPT-4o-mini	0.04	80.20	88.70
GPT-4o	3.31	38.85	70.95
USP-8b	77.73	98.05	97.55
UserLM-1b	78.96	91.30	93.55
UserLM-8b	80.21	93.95	94.65
Human (estimate)	90.15	—	—

Evaluating Simulator Robustness



Naturalness

Do simulator utterances resemble natural user text?

User Simulator



(↑)



(↑)



(↑)

Llama3.2-1b-Instruct

0.14

77.55

54.95

User LMs are more robust simulators than assistant LMs
(stick to their user role, do not change their intent easily)

USP-8b

77.73

98.05

97.55

UserLM-1b

78.96

91.30

93.55

UserLM-8b

80.21

93.95

94.65

Human (estimate)

90.15

—

—



User Intent Adherence

Does the simulator stick to its user intent?

Flipping The Dialogue

1

Training User LMs

How do we train a user language model?

2

Intrinsic Evaluation

How do we evaluate a user language model?

3

Extrinsic Evaluation

How do user language models help in practice?

Extrinsic Eval Setup: Multi-Turn Evaluation of Assistants

We simulate multi-turn conversations between a GPT-4o assistant and UserLM-8b

Tasks: 65 task intents that involve users completing:

- Math word problems (*based on GSM8k* (Cobbe et al., 2021))
- Writing Python programs (*based on HumanEval* (Chen et al., 2021))

Baselines: We compare to simulators based on prompted GPT-4o and 4o-mini

Evaluation: We analyze various properties of the simulated conversations for each simulator and compare to how the performance of the assistant changes

User Simulator

4o-mini	GPT-4o	UserLM-8b
---------	--------	-----------

Intent Coverage

Intent Coverage (%)	86.6	84.7	76.7
---------------------	------	------	------

User Simulator

4o-mini	GPT-4o	UserLM-8b
---------	--------	-----------

Intent Coverage

Intent Coverage (%)	86.6	84.7	76.7
---------------------	------	------	------

Information Diversity

%Repeat Required	31.8	9.4	54.3
%Skip Non-Required	10.9	14.6	37.7
%Add Demands	9.5	1.1	43.8

UserLM-8b introduces more nuances to the conversation

User Simulator

4o-mini	GPT-4o	UserLM-8b
---------	--------	-----------

Intent Coverage

Intent Coverage (%)	86.6	84.7	76.7
---------------------	------	------	------

Information Diversity

%Repeat Required	31.8	9.4	54.3
%Skip Non-Required	10.9	14.6	37.7
%Add Demands	9.5	1.1	43.8

Pace Diversity

Turn Variance	0.9	0.6	2.8
Turn Range	3.7-5.7	4.0-5.4	2.1-6.7

UserLM-8b introduces more nuances to the conversation

UserLM-8b simulates more varied conversational paces

User Simulator

4o-mini	GPT-4o	UserLM-8b
---------	--------	-----------

Intent Coverage

Intent Coverage (%)	86.6	84.7	76.7
---------------------	------	------	------

Information Diversity

%Repeat Required	31.8	9.4	54.3
%Skip Non-Required	10.9	14.6	37.7
%Add Demands	9.5	1.1	43.8

Pace Diversity

Turn Variance	0.9	0.6	2.8
Turn Range	3.7-5.7	4.0-5.4	2.1-6.7

Lexical Diversity

Unigram Difference	0.43	0.40	0.71
--------------------	------	------	------

UserLM-8b introduces more nuances to the conversation

UserLM-8b simulates more varied conversational paces

UserLM-8b uses more varied phrasing of requests

User Simulator

	4o-mini	GPT-4o	UserLM-8b
Intent Coverage			
Intent Coverage (%)	86.6	84.7	76.7
Information Diversity			
%Repeat Required	31.8	9.4	54.3
%Skip Non-Required	10.9	14.6	37.7
%Add Demands	9.5	1.1	43.8
Pace Diversity			
Turn Variance	0.9	0.6	2.8
Turn Range	3.7-5.7	4.0-5.4	2.1-6.7
Lexical Diversity			
Unigram Difference	0.43	0.40	0.71
Assistant Diversity			
Assistant Score	73.2	74.6	57.4

UserLM-8b introduces more nuances to the conversation

UserLM-8b simulates more varied conversational paces

UserLM-8b uses more varied phrasing of requests

Better estimate of assistant performance

User Simulator

	4o-mini	GPT-4o	UserLM-8b
--	---------	--------	-----------

Intent Coverage

Intent Coverage (%)	86.6	84.7	76.7
---------------------	------	------	------

Information Diversity

Assistants struggle to handle user nuances in multi-turn settings
(UserLM-8b covers the necessary information to solve the tasks, but the nuances it introduces causes assistant failure)

Turn Range	3.7-5.7	4.0-5.4	2.1-6.7
------------	---------	---------	---------

Lexical Diversity

Unigram Difference	0.43	0.40	0.71
--------------------	------	------	------

Assistant Diversity

Assistant Score	73.2	74.6	57.4
-----------------	------	------	------

UserLM-8b simulates more varied conversational paces

UserLM-8b uses more varied phrasing of requests

Better estimate of assistant performance

Where do we go from here?

Takeaways:

- The assistant role models are post-trained to play limits their capabilities for user simulation
- Language models should be trained to play distinct and opposing roles: ***user*** and ***assistant***
- User LMs will enable the development of robust assistant that are ready for real-world nuances

Notes for the future:

- More releases of better and bigger base models will be needed
- User LMs can serve purposes beyond multi-turn eval (synthetic data, user-centered judge, etc.)
- More personalized user LMs will be needed to simulate specific sub-populations

ආචාර්ය شُكْرًا Merci 谢谢 धन्यवाद Asante Teşekkürler
ありがとう Gracias متشكرم நன்றி Obrigado Thank You

microsoft/**UserLM-8b** like 365 Follow Microsoft 19k

Text Generation Safetensors allenai/WildChat-1M English

Downloads last month
1,079



microsoft/UserLM-8b model card

Model description

Unlike typical LLMs that are trained to play the role of the "assistant" in conversation, we trained UserLM-8b to simulate the "user" role in conversation (by training it to predict user turns in a large corpus of conversations called WildChat). This model is useful in simulating more realistic conversations, which is in turn useful in the development of more robust assistants.

Safetensors

Model size 8B params Tensor type F32 Chat template

Files info

UserLM-8b is available at: 🙌 <https://huggingface.co/microsoft/UserLM-8b>

Feel free to follow up with me on 🐦 @tareknaous